

ASSIGNMENT-2

- 1. In logistic regression, what is the logistic function (sigmoid function) and how is it used to compute probabilities?**

Ans: Logistic Function (Sigmoid Function) in Logistic Regression:

The logistic function, also known as the sigmoid function, is a mathematical function that maps any real-valued number to a value between 0 and 1. It has an S-shaped curve. In logistic regression, the logistic function is used to model the probability that a given input belongs to a certain class. The output of the logistic function represents the probability of the input belonging to the positive class (class 1).

- 2. When constructing a decision tree, what criterion is commonly used to split nodes, and how is it calculated?**

Ans: Criterion for Splitting Nodes in Decision Trees:

The commonly used criterion for splitting nodes in decision trees includes measures like Gini impurity and information gain (entropy).

Gini impurity measures the probability of incorrectly classifying a randomly chosen element if it was randomly labelled according to the distribution of labels in the node.

Information gain, based on entropy, measures the reduction in entropy (or uncertainty) of the target variable after the split.

- 3. Explain the concept of entropy and information gain in the context of decision tree construction.**

Ans: Entropy and Information Gain:

Entropy is a measure of randomness or uncertainty in a dataset. In decision tree construction, entropy is used to calculate the homogeneity of a sample.

Information gain measures the reduction in entropy or uncertainty after a dataset is split based on a feature. It helps in deciding the best feature to split the data on at each node of the tree.

- 4. How does the random forest algorithm utilize bagging and feature randomization to improve classification accuracy?**

Ans: Random Forest Algorithm and Bagging:

Random Forest utilizes bagging (bootstrap aggregating) by training multiple decision trees on random subsets of the training data and then combining their predictions to reduce overfitting and improve generalization.

Feature randomization is used by randomly selecting a subset of features at each split in each decision tree to increase diversity among the trees and prevent overfitting.

- 5. What distance metric is typically used in k-nearest neighbours (KNN) classification, and how does it impact the algorithm's performance?**

Ans: Distance Metric in K-Nearest Neighbours (KNN):

The Euclidean distance metric is typically used in KNN classification, although other distance metrics such as Manhattan distance, Minkowski distance, etc., can also be used.

The choice of distance metric impacts how the algorithm measures similarity between data points, which in turn affects its performance.

6. Describe the Naïve-Bayes assumption of feature independence and its implications for classification.

Ans: Naïve-Bayes Assumption of Feature Independence:

Naïve-Bayes assumes that the features are conditionally independent given the class label. This means that the presence of a particular feature in a class is independent of the presence of other features.

Despite this simplifying assumption, Naïve-Bayes often performs well in practice and is computationally efficient.

7. In SVMs, what is the role of the kernel function, and what are some commonly used kernel functions?

Ans: Role of Kernel Function in SVMs:

The kernel function in SVMs is used to transform the input data into a higher-dimensional space where it becomes linearly separable.

Commonly used kernel functions include linear kernel, polynomial kernel, Gaussian (RBF) kernel, and sigmoid kernel.

8. Discuss the bias-variance trade-off in the context of model complexity and overfitting.

Ans: Bias-Variance Tradeoff:

The bias-variance tradeoff refers to the balance between bias (error due to overly simplistic assumptions) and variance (error due to sensitivity to fluctuations in the training set) in machine learning models.

Increasing model complexity typically reduces bias but increases variance, and vice versa. Overfitting occurs when the model captures noise in the training data instead of the underlying pattern.

9. How does TensorFlow facilitate the creation and training of neural networks?

Ans: TensorFlow for Neural Networks:

TensorFlow facilitates the creation and training of neural networks by providing a flexible framework for building computational graphs, automatic differentiation for optimizing model parameters, GPU acceleration for faster computations, and high-level APIs like Keras for building and training neural networks more easily.

10. Explain the concept of cross-validation and its importance in evaluating model performance.

Ans: Cross-Validation:

Cross-validation is a technique used to assess the generalization performance of a predictive model. It involves partitioning the dataset into multiple subsets, training the model on some subsets, and evaluating it on the remaining subset.

Cross-validation helps in estimating how well the model will perform on unseen data and reduces the risk of overfitting.

11. What techniques can be employed to handle overfitting in machine learning models?

Ans: Techniques for Handling Overfitting:

Regularization techniques like L1 and L2 regularization penalize large model coefficients to prevent overfitting. Feature selection to reduce the complexity of the model and focus on relevant features.

Early stopping during training to prevent the model from learning noise in the data.

12. What is the purpose of regularization in machine learning, and how does it work?

Ans: Regularization in Machine Learning:

Regularization is a technique used to prevent overfitting by adding a penalty term to the model's loss function, which discourages overly complex models.

Common regularization techniques include L1 regularization (lasso), L2 regularization (ridge), and elastic net regularization.

13. Describe the role of hyper-parameters in machine learning models and how they are tuned for optimal performance.

Ans: Hyperparameters in Machine Learning Models:

Hyperparameters are parameters that are set prior to training and control the learning process of the model. They are tuned to optimize the model's performance, often through techniques like grid search, random search, or Bayesian optimization.

14. What are precision and recall, and how do they differ from accuracy in classification evaluation?

Ans: Precision, Recall, and Accuracy:

Precision measures the proportion of true positive predictions among all positive predictions.

Recall measures the proportion of true positive predictions among all actual positives.

Accuracy measures the proportion of correct predictions among all predictions.

15. Explain the ROC curve and how it is used to visualize the performance of binary classifiers.

Ans: ROC Curve:

The Receiver Operating Characteristic (ROC) curve is a graphical plot that illustrates the performance of a binary classifier across different threshold settings.

It plots the true positive rate (TPR) against the false positive rate (FPR) at various threshold values.

The area under the ROC curve (AUC) is used as a summary measure of the classifier's performance, with a higher AUC indicating better performance.